

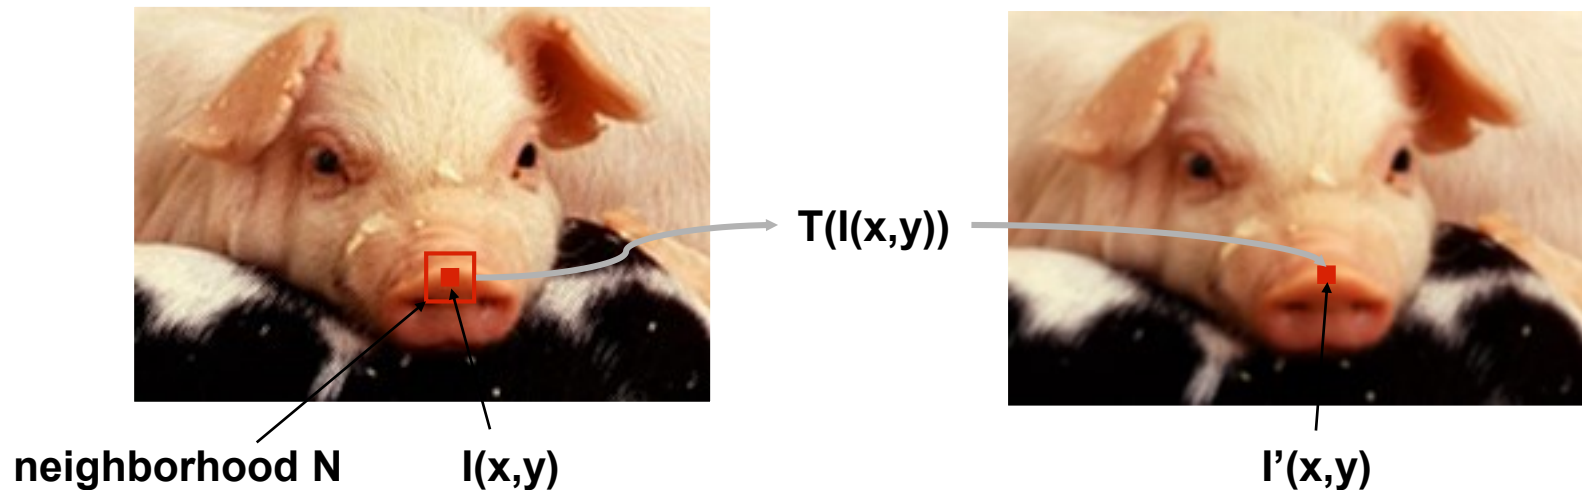
# IMAGE PROCESSING

- Image Enhancement
  - Brightness mapping
  - Contrast stretching/enhancement
  - Histogram modification
  - Noise Reduction
  - .....
- Mathematical Techniques
  - Convolution
    - Mean filtering
    - Gaussian filtering
- Edge and Line Detection and Extraction
- Contour Extraction
- Corner Detection
- Region Segmentation

- Thresholding
  - threshold selection (manual & automatic)
- Transformations for contrast enhancement
  - linear
    - linear stretching
  - non-linear
    - power law
    - logarithmic
    - equalization
    - CLAHE

- Linear filters: mean and Gaussian
  - convolution operation
- Non-linear filters
  - median
  - anisotropic diffusion filter
  - bilateral filter
  - ...
- Frequency domain filters

- Goal: improve the 'visual quality' of the image
  - for human viewing
  - for subsequent processing
- Two typical methods
  - spatial domain techniques....
    - operate directly on image pixels
  - frequency domain techniques....
    - operate on the Fourier transform of the image
- No general theory of 'visual quality'
  - General assumption: if it looks better, it is better
  - Often not a good assumption



- Transformation T
  - point - pixel to pixel
  - local - local area to pixel
  - global - output value at a specific coordinate depends on all values in the input image. (ex: DFT)
- Local - neighborhoods
  - typically quadrangular
  - typically an odd size: 3x3, 5x5, etc. (why odd size? see later)
  - centered on pixel  $I(x,y)$

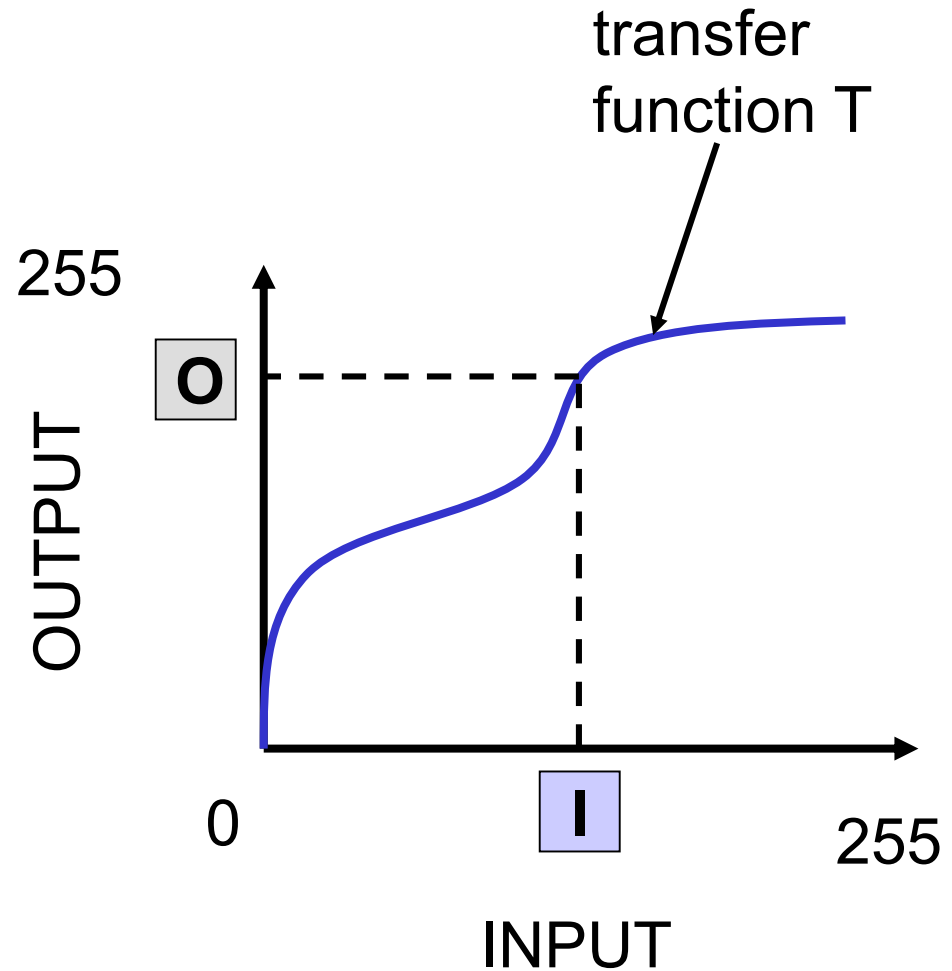
$$I'(x,y) = T(I(x,y))$$

# Point transformations

Histogram-based transformations

$$O = T(I)$$


Input pixel value,  $I$ , mapped to output pixel value,  $O$ , via transfer function  $T$ .





- T is a point-to-point transformation
  - only information at  $I(x,y)$  used to generate  $I'(x,y)$
- Thresholding

$$I'(x,y) = \begin{cases} I_{\max} & \text{if } I(x,y) > t \\ I_{\min} & \text{if } I(x,y) \leq t \end{cases}$$



Color image




Graylevel image













Thresholded graylevel image ( $t=89$ )

- Arbitrary selection
  - select visually
- Use image histogram



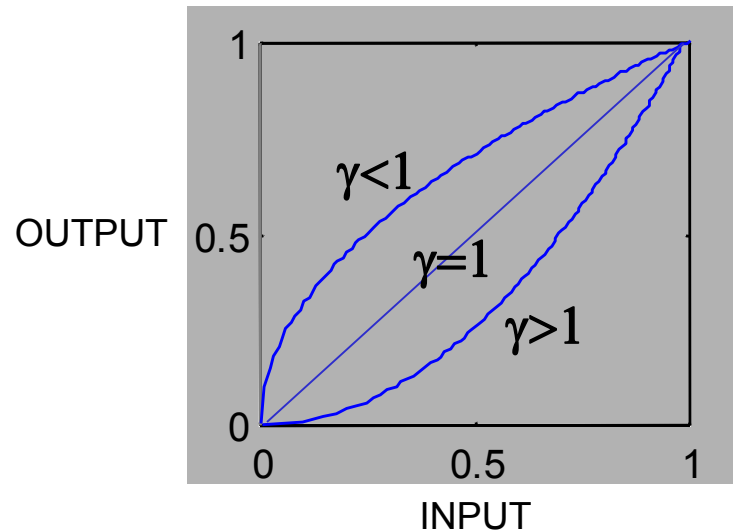
Later, we'll be back to threshold selection methods

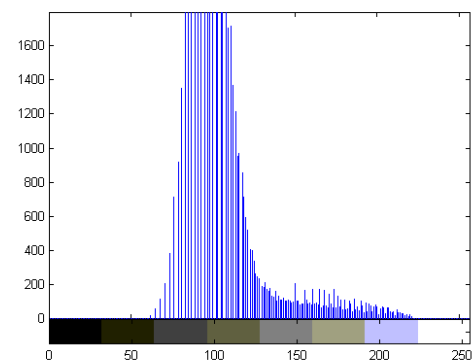


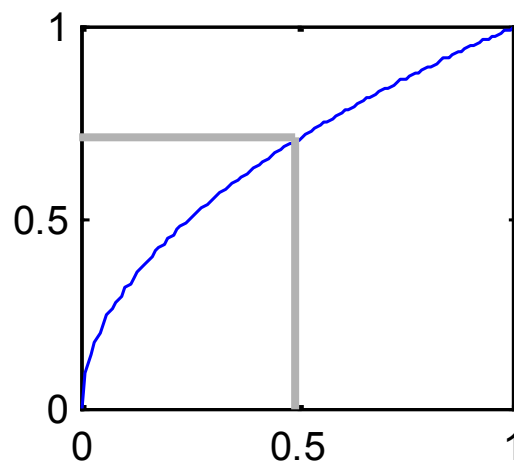
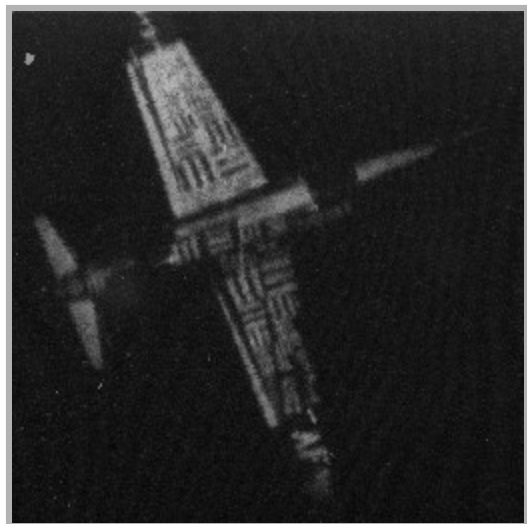


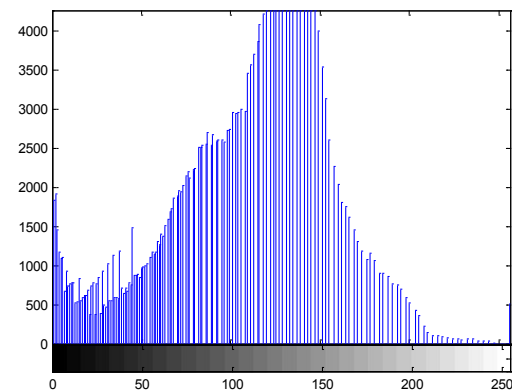
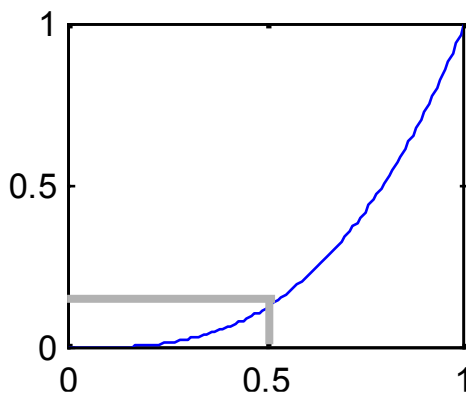
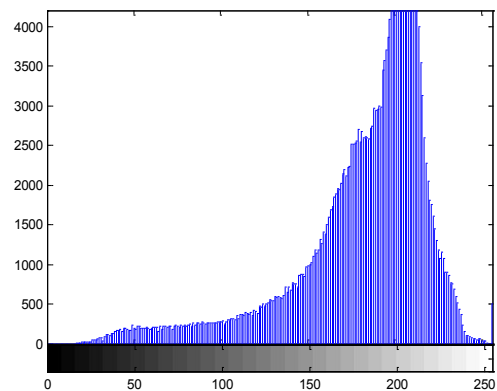
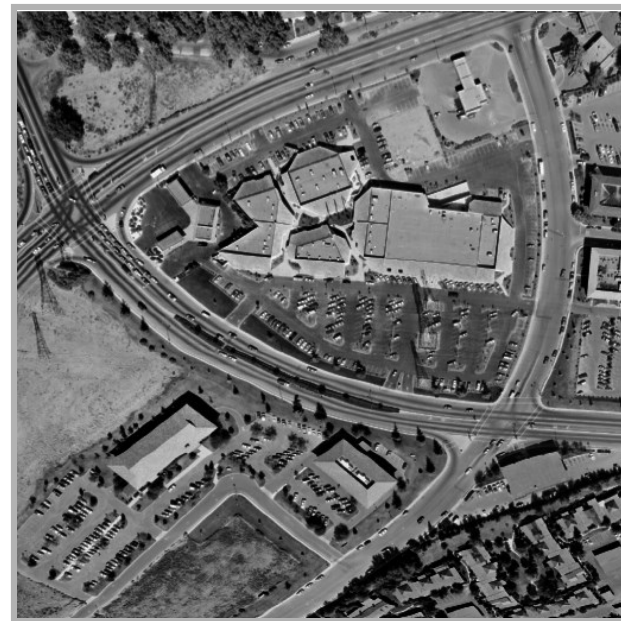


$$O = I^{\gamma}$$

- $\gamma < 1$  to enhance contrast in dark regions
- $\gamma > 1$  to enhance contrast in bright regions.





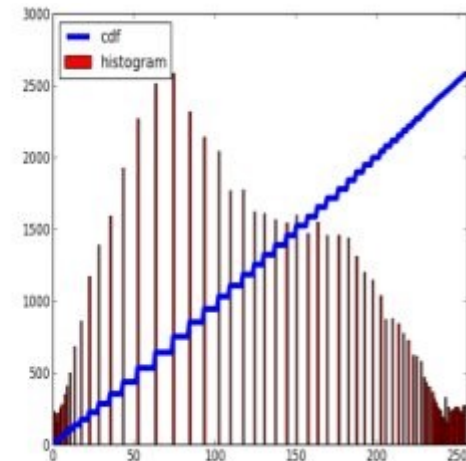
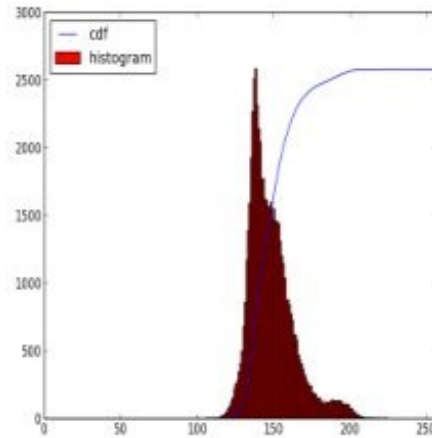


- Technique can be applied to color images
  - same curve to all color bands
  - different curves to separate color bands:



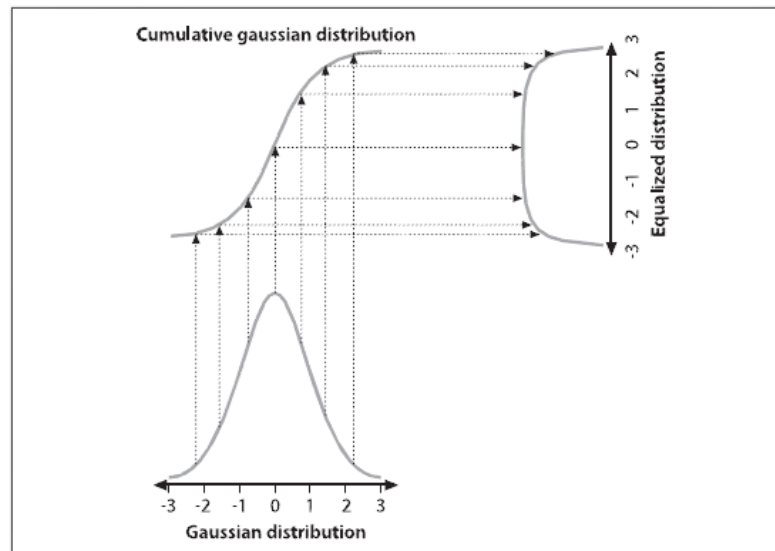
- but ... **be careful**, colors may become distorted...!
- Point transformations are usually applied using **LUT's (Look Up Tables)**





source: <http://docs.opencv.org/>

After equalization  
the cumulative  
distribution function  
is almost linear



Using the cumulative distribution function  
to equalize a Gaussian distribution



Original



Histogram  
stretching




Histogram  
equalization

- The contrast enhancement is better after the histogram equalization, which more easily detects structures located in the shade.
- In fact any strongly represented gray-level is stretched while any weakly represented graylevel is merged with other close levels


Original



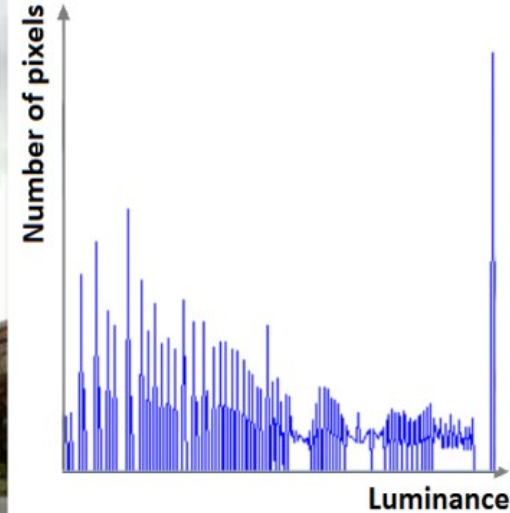
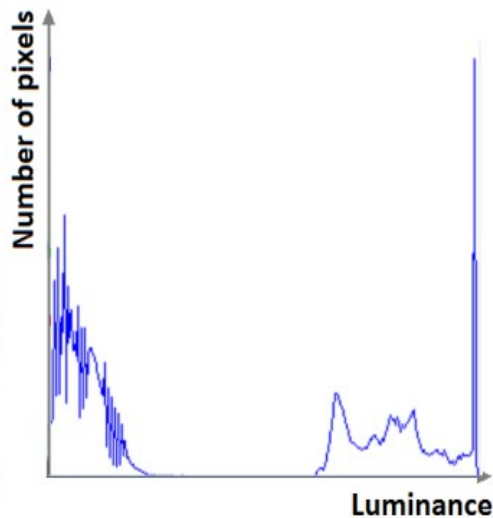
$\gamma > 1$



Histogram  
equalization



- In color images, only the luminance channel is usually equalized as otherwise the colors can become distorted.

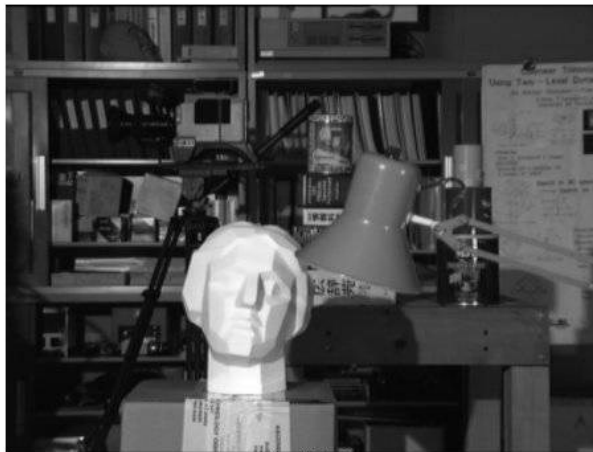


- Acquisition process degrades image
- Brightness and contrast enhancement implemented by pixel operations
- No one algorithm universally useful
- $\gamma > 1$  enhances contrast in bright images
- $\gamma < 1$  enhances contrast in dark images
- Transfer function for histogram equalization proportional to cumulative histogram
- It is essential a process of trial-and-error to determine whether a particular type of images will benefit from histogram transformation operations.

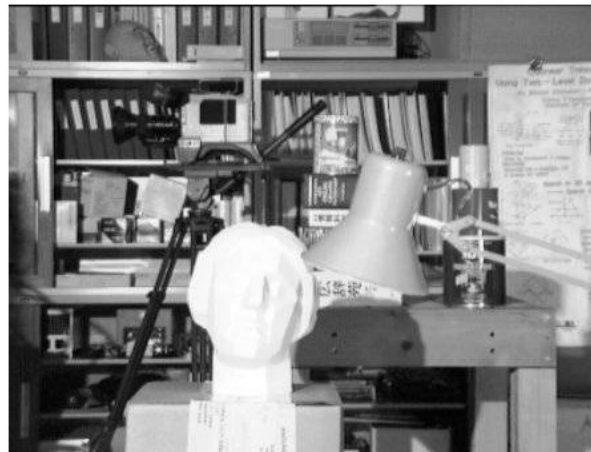


- **CLAHE – Contrast Limited Adaptive Histogram Equalization**

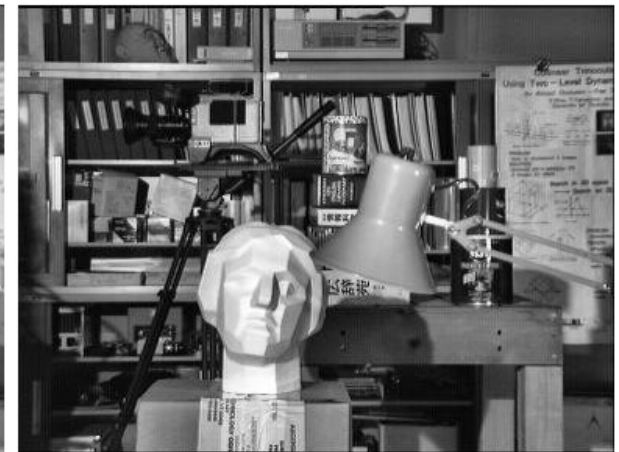
- Considering the global contrast of the image is not a good idea, in some cases ... (look at the face of the statue, in the middle image)
- For some images, it might be preferable to apply different kinds of equalization in different regions.
- Instead of computing a single curve, the image is divided into  $M \times M$  pixel non-overlapped sub-blocks and separate histogram equalization is performed in each sub-block.
- To avoid blocking artifacts (i.e., intensity discontinuities at block boundaries) in the resulting image, the equalization functions are smoothly interpolated as we move between blocks.
- This technique is known as adaptive histogram equalization (AHE) and its contrast limited (gain-limited, to avoid noise amplification) version is known as CLAHE.



Original image



After global equalization  
(notice the effect on the face of the statue)



After CLAHE

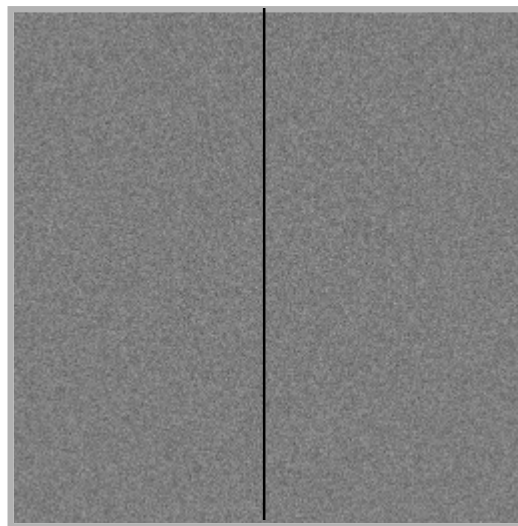
# Filtering

- What is noise?
- How is noise reduction performed?
  - Noise reduction from first principles
  - Neighbourhood operators
    - linear filters (low pass)
    - non-linear filters (median)



image

+



noise


=



'grainy' image

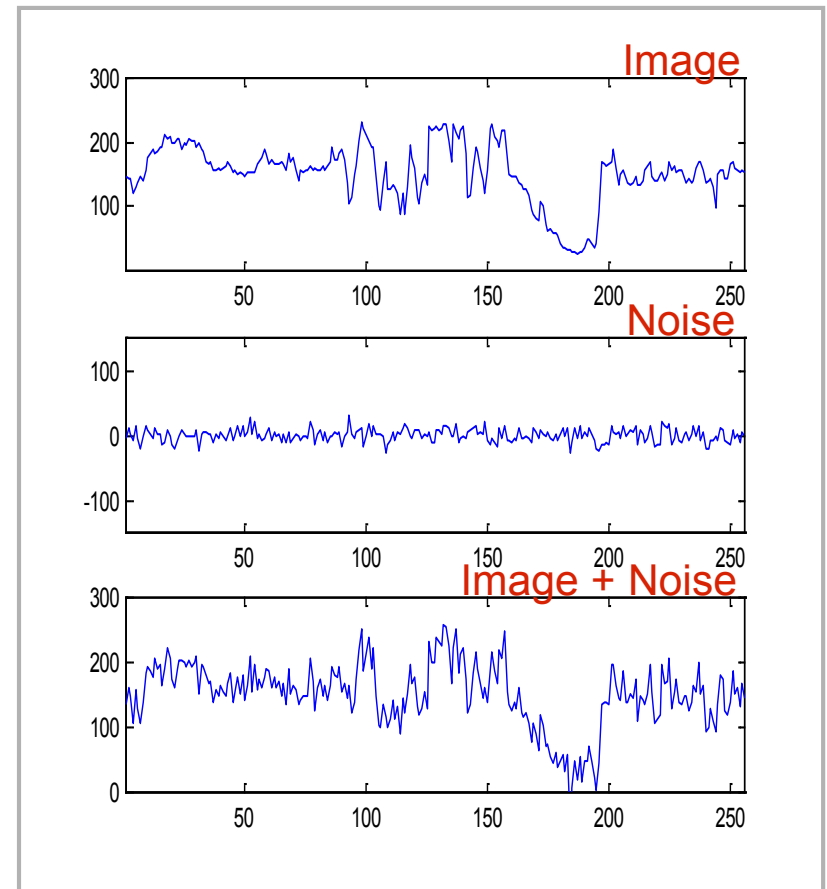


- Sources of noise = CCD chip.
- Electronic signal fluctuations in detector.
  - Caused by thermal energy.
  - Worse for infra-red sensors.
- Other electronics
- Transmission (analog)

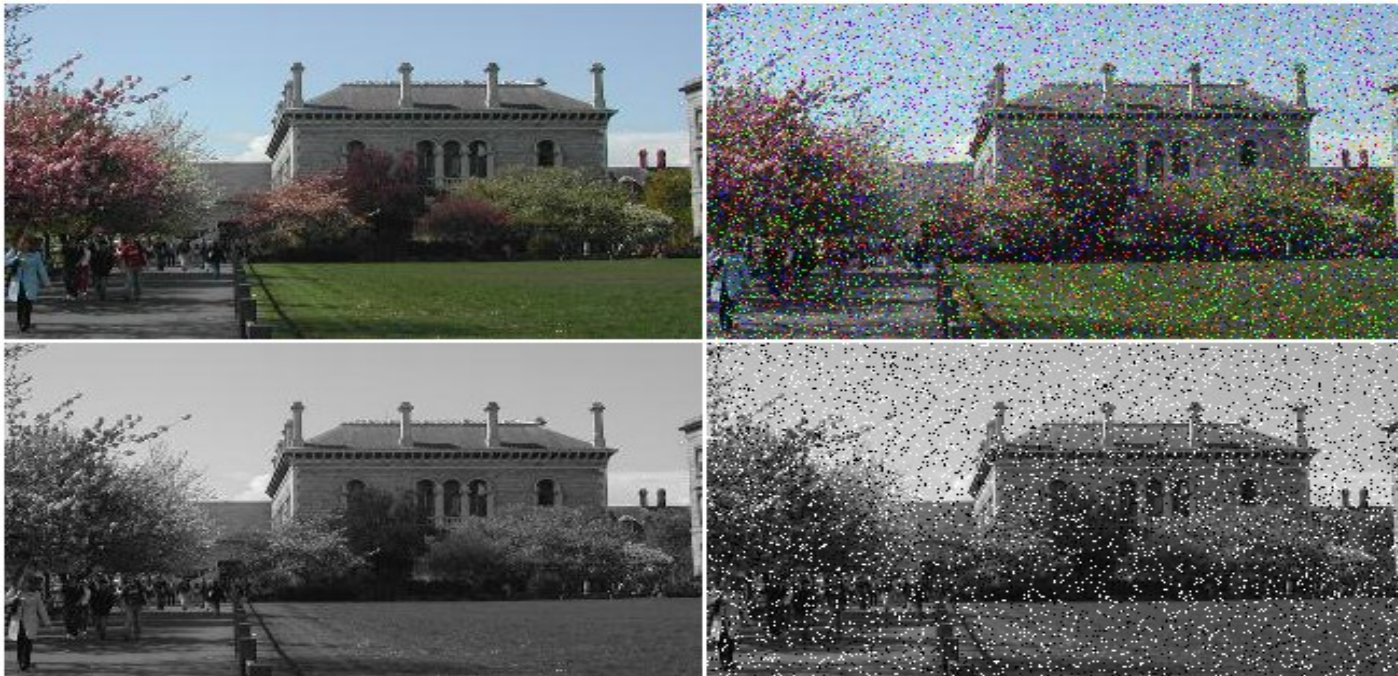


Radiation from the long wavelength IR band is used in most infrared imaging applications

- Plot of image brightness.
  - Noise models:
    - additive
      - $f(i,j) = g(i,j) + v(i,j)$ 
        - Gaussian
        - salt and pepper
    - multiplicative
      - $f(i,j) = g(i,j) + g(i,j).v(i,j)$
- where
- $f(i,j)$  - acquired signal (noisy)
  - $g(i,j)$  - uncorrupted signal
  - $v(i,j)$  - noise component
- Noise fluctuations are rapid
    - high frequency.



- Impulse noise
  - Noise is maximum or minimum values

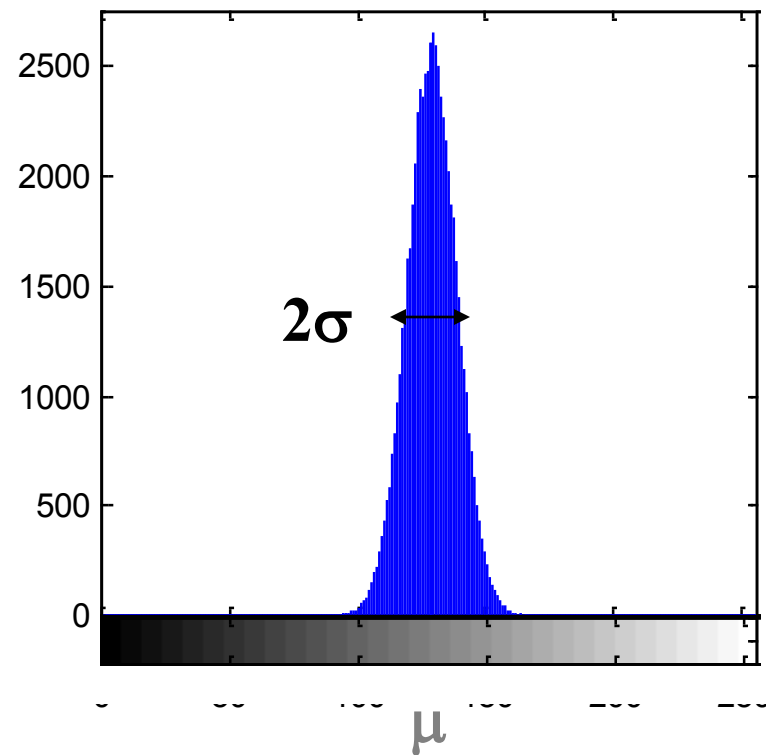


- Good approximation to real noise
- Distribution is Gaussian (mean & standard deviation)



- Plot noise histogram
- Typical noise distribution is normal or Gaussian
- Mean(noise)  $\mu = 0$
- Standard deviation  $\sigma$

$$\eta(x) = \frac{1}{\sqrt{2\pi}\sigma} \exp\left[-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}\right]$$



- Noise varies above and below uncorrupted image.

Image

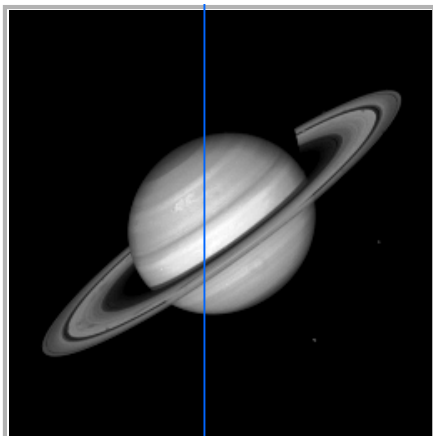
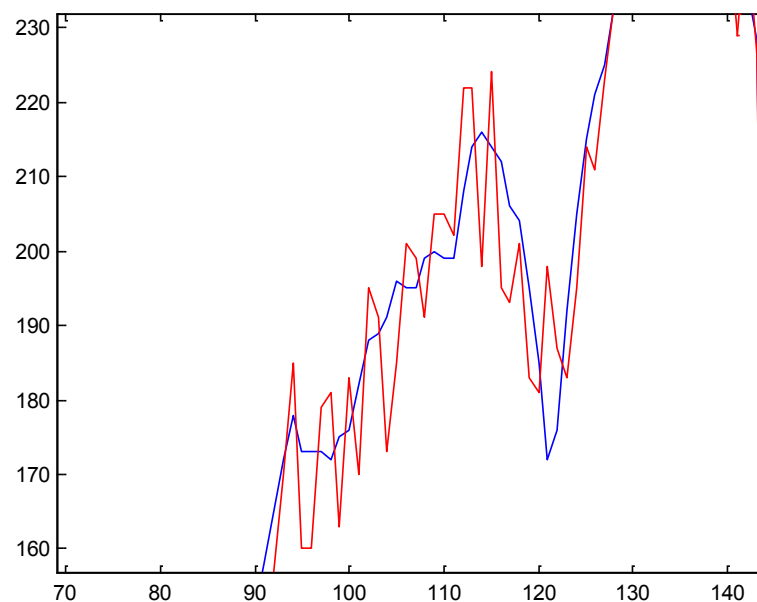



Image  
+  
Noise







- Removing (?) or reducing noise...
- Linear smoothing transformations
  - frame averaging (=> acquire several frames of static scene)
  - local averaging
  - Gaussian smoothing
- Non-linear transformations
  - median filter
  - rotating mask

Image Data:


10	12	40	16	19	10
14	22	52	10	55	41
10	14	51	21	14	10
32	22	9	<b>9</b>	19	14
41	18	9	22	27	11
10	7	8	8	4	5

Mask / Filter / Kernel:

1	1	1	1
0	1	1	1
-1	1	1	1
	-1	0	1

$$F \circ I(x, y) = \sum_{j=-N}^N \sum_{i=-N}^N \overset{\text{Filter}}{F(i, j)} \overset{\text{Image}}{I(x+i, y+j)}$$





- Kernel is aligned with pixel in image, multiplicative sum is computed, normalized,
- and stored in result image.
- Process is repeated across image.
- What happens when kernel is near edge of input image?

$$\left[ \begin{array}{l} 1*51 + 1*21 + 1*14 + \\ 1*8 + 1*9 + 1*19 + \\ 1*19 + 1*22 + 1*27 \end{array} \right] \left( \frac{1}{9} \right) = 21$$


1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

*F*

← note the factor

- missing samples are zero
- missing samples are gray
- copying last lines
- reflected indexing (mirror)
- circular indexing (periodic)
- reduce size of resulting image

- OpenCV  
allows some control on how to do this;  
see `CopyMakeBorder()`: copies the source  
2D array into the interior of the destination  
array and makes a border of the specified  
type around the copied area



1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

***F***

8	3	4	5
7	6	4	5
4	5	7	8
6	5	5	6

***I***

8	8	3	4	5	5
8	8	3	4	5	5
7	7	6	4	5	5
4	4	5	7	8	8
6	6	5	5	6	6
6	6	5	5	6	6

***I with padded boundaries***

6.44	5.22	4.33	4.67
5.78	5.33	5.22	5.67
5.56	5.44	5.67	6.00
5.22	5.33	5.78	6.33

***J = FoI***

- A **convolution** operation is a correlation where the filter is flipped both horizontally and vertically before being applied to the image:

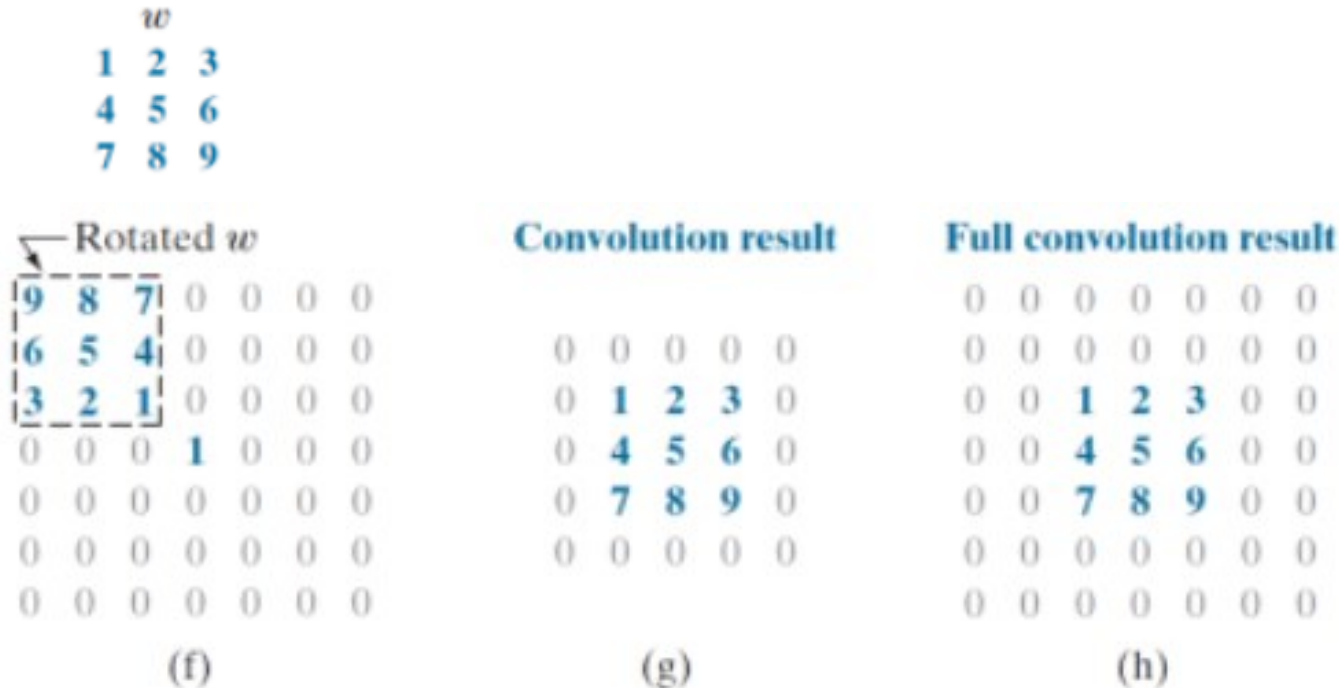
$$F * I(x, y) = \sum_{j=-N}^N \sum_{i=-N}^N F(i, j) I(x - i, y - j)$$

- Suppose  $F$  is a Gaussian or mean kernel.  
How does convolution differ from cross-correlation?

- Extremely important concept in computer vision, image processing, signal processing, etc.
- Lots of related mathematics (**we won't do**)
- General idea: reduce a filtering operation to the repeated application of a mask (or filter kernel) to the image
  - Kernel can be thought of as an  $N \times N$  image
  - $N$  is usually odd so kernel has a central pixel
- In practice
  - (flip kernel)
  - Align kernel center pixel with an image pixel
  - Pointwise multiply each kernel pixel value with corresponding image pixel value and add results
  - Resulting sum is normalized by kernel weight
  - Result is the value of the pixel centered on the kernel



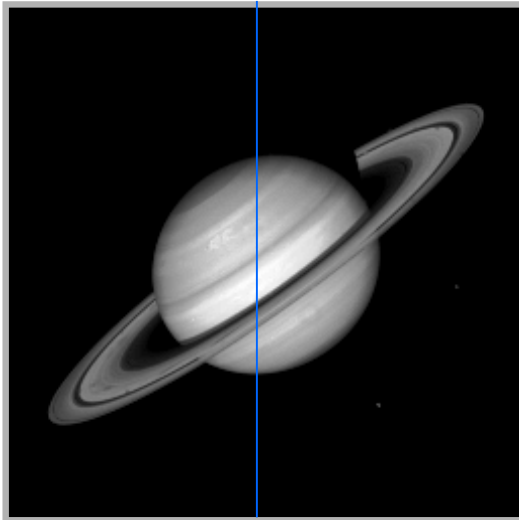
## Convolution





source: <https://towardsdatascience.com/convolution-vs-correlation-af868b6b4fb5>

## Notes:

- 1) Correlation and convolution are identical when the filter is symmetric.
- 2) The key difference between the two is that convolution is associative, that is, if  $F$  and  $G$  are filters, then  $F*(G*Img) = (F*G)*Img$
- 3) In general, people use **convolution** for image processing operations such as smoothing, and they use **correlation** to match a template to an image (*see later*).



Uncorrupted Image



Uncorrupted Image + Noise




Image + Noise - Blurred

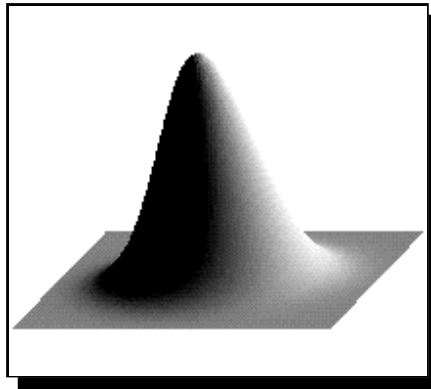


- Technique relies on high frequency noise fluctuations being 'blocked' by filter.  
Hence, **low-pass** filter.
- Fine detail in image may also be smoothed.
- Balance between  
keeping image fine detail and reducing noise.

- Saturn image  
(*previous slide*)  
has coarse detail
- Boat image  
contains fine detail
- Noise reduced but  
fine detail also smoothed



- Smoothing operator should be
  - 'tunable' in what it leaves behind
  - smooth and localized in image space.
- One operator which satisfies these two constraints is the Gaussian




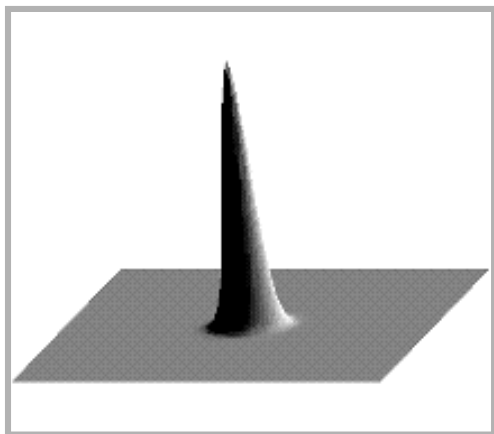
- *OpenCV: Smooth()*

- The two-dimensional Gaussian distribution is defined by:


$$G(x,y) = \frac{1}{\sigma \sqrt{2\pi}} e^{-\left[ \frac{(x^2 + y^2)}{2 \sigma^2} \right]}$$

- From this distribution, can generate smoothing masks whose width depends upon  $\sigma$ :








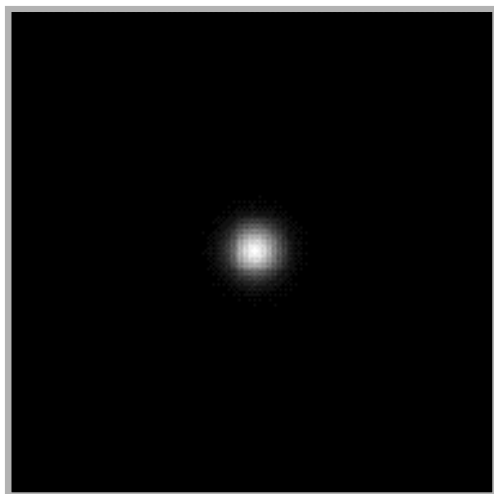
$$\sigma^2 = .25$$



$$\sigma^2 = 1.0$$



$$\sigma^2 = 4.0$$



- Choose  $\sigma^2 = 2$  and  $n = 7$ , then:


		j						
		-3	-2	-1	0	1	2	3
i	-3	.011	.039	.082	.105	.082	.039	.011
	-2	.039	.135	.287	.368	.287	.135	.039
	-1	.082	.287	.606	.779	.606	.287	.082
	0	.105	.039	.779	1.000	.779	.368	.105
	1	.082	.287	.606	.779	.606	.287	.082
	2	.039	.135	.287	.368	.287	.135	.039
	3	.011	.039	.082	.105	.082	.039	.011

1	4	7	10	7	4	1
4	12	26	33	26	12	4
7	26	55	71	55	26	7
10	33	71	91	71	33	10
7	26	55	71	55	26	7
4	12	26	33	26	12	4
1	4	7	10	7	4	1

7x7 Gaussian filter

$$\frac{W(1,2)}{k} = \exp\left(-\frac{1^2 + 2^2}{2 \cdot 2}\right)$$

To make this value 1,  
choose  $k=91$





7x7 Gaussian kernel



15x15 Gaussian kernel

- Gaussian is not the only choice, but it has a number of important properties
  - If we convolve a Gaussian with another Gaussian, the result is a Gaussian
    - This is called linear scale space

$$G_{\sigma_1} * G_{\sigma_2} = G_{\sqrt{\sigma_1^2 + \sigma_2^2}}.$$

- Efficiency: separable

$$\begin{aligned} G_{\sigma}(x, y) &= \frac{1}{2\pi\sigma^2} \exp\left(-\frac{(x^2 + y^2)}{2\sigma^2}\right) \\ &= \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(x^2)}{2\sigma^2}\right)\right) \times \left(\frac{1}{\sqrt{2\pi}\sigma} \exp\left(-\frac{(y^2)}{2\sigma^2}\right)\right), \end{aligned}$$

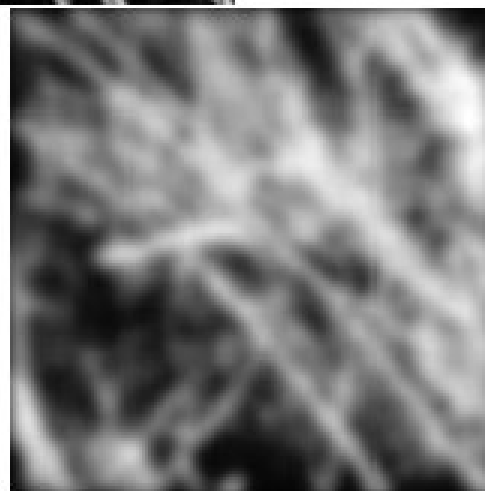




Original image



After mean filtering



After Gaussian filtering

Image




Blurred image



-

=



- Compute median of points in neighborhood
- Nonlinear filter
  - example of a larger class of filters named "rank-filters" (ex: min, median, max)
- Has a tendency to blur detail less than averaging
- Works very well for 'shot' or 'salt and pepper' noise

Linear vs Non-linear

$$\text{Mean}(I1+I2) = \text{Mean}(I1) + \text{Mean}(I2)$$

$$\text{Median}(I1+I2) \neq \text{Median}(I1) + \text{Median}(I2)$$



Original



Low-pass



Median

- Mean
  - Linear
  - Signal frequencies shared with noise are lost, resulting in blurring.
  - Impulsive noise is diffused but not removed.
  - It spreads the noise, resulting in blurring.
  - Blurs edges.
- Median
  - Non-linear
  - Does not spread the noise.
  - Can remove spike noise.
  - Preserves (some) edges.
  - Small details (small regions, thin edges, ...) may be lost.
  - Expensive to run.

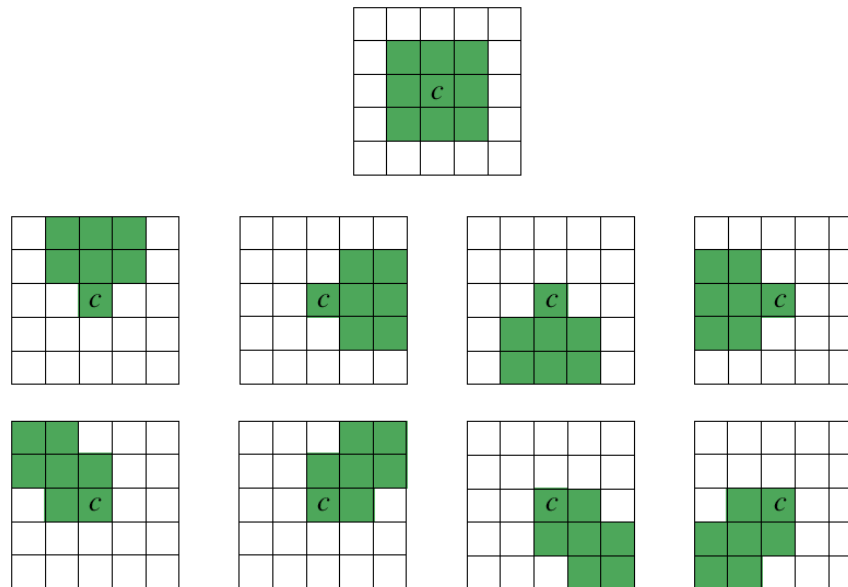
Smooth (average, blur) an image without disturbing

- sharpness or
- position

of the edges

- Nagao-Matsuyama filter
- Kuwahara filter
- Anisotropic diffusion filters (Perona & Malik, .....)
- Bilateral filtering
- ...

- Calculate the variance within nine subwindows of a 5x5 moving window
- Output value is the **mean** of the **subwindow with the smallest variance**
- Nine subwindows used:



- Principle:
  - divide filter mask into four regions (a, b, c, d).
  - in each compute the mean brightness and the variance
  - the output value of the center pixel (abcd) in the window is the **mean** value of that **region** that has the **smallest variance**.

a	a	ab	b	b
a	a	ab	b	b
ac	ac	abcd	bd	bd
c	c	cd	d	d
c	c	cd	d	d



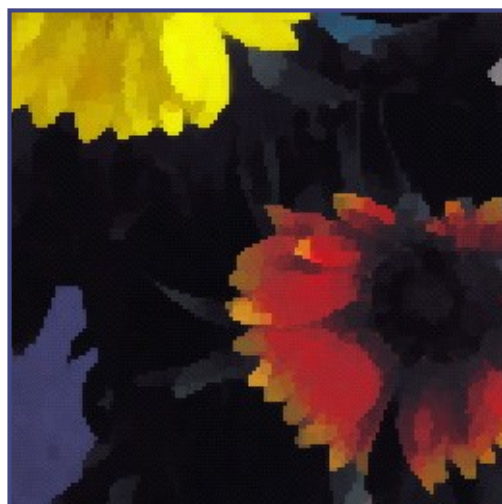




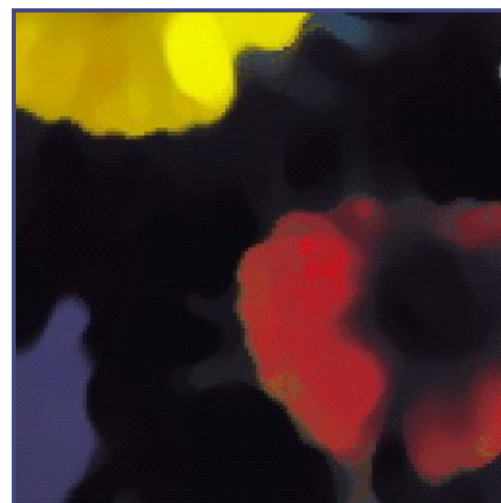
**Original**



**Median (1 iteration)**



**Kuwahara**



**Median (10 iterations)**



- The anisotropic diffusion is a nonlinear filter that smoothes the intraregions of an image without blurring the strong edges
  - Encourages the **smoothing in homogeneous regions in preference to** smoothing across the boundaries.
- Based on the solution of a partial differential equation, inspired in heat diffusion equation
  - *The algorithm results from the discretization of the non-linear diffusion equation.*
  - *The derivatives are approximated by differences.*
- Diffusion methods average over extended regions by solving partial differential equations, and are therefore inherently iterative. Iteration may raise issues of stability and, depending on the computational architecture, efficiency.
- This algorithm is also used to detect the edges



Original



Perona, Malik, 1987

$$\frac{\partial u}{\partial t} = \operatorname{div} \left( \frac{Du}{|Du|^2 + \lambda^2} \right)$$

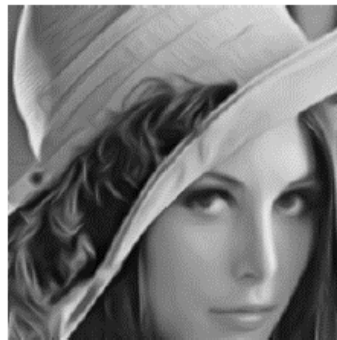
Rudin, Osher,  
Fatemi, 1992

$$\frac{\partial u}{\partial t} = \operatorname{div} \left( \frac{Du}{|Du|} \right)$$



Alvarez, Lions, 1992

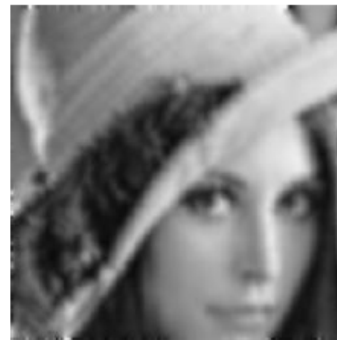
$$\frac{\partial u}{\partial t} = \frac{|Du|}{|k * Du|} \operatorname{div} \left( \frac{Du}{|Du|} \right)$$



Weickert, 1994

$$\frac{\partial u}{\partial t} = D^2 u(d, d)$$

$$d = \operatorname{SEigen}(k * (Du \otimes Du))$$



Caselles, Sbert, 1997

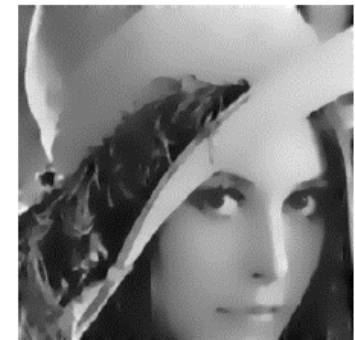
$$\frac{\partial u}{\partial t} = \frac{1}{|Du|^2} D^2 u(Du, Du)$$



Zhong Carmona, 1998



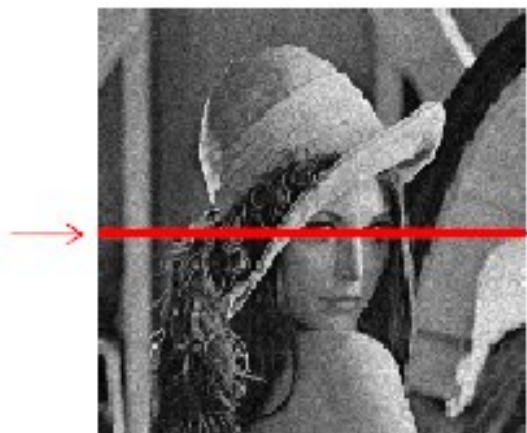
$$\frac{\partial u}{\partial t} = D^2 u(d, d)$$

$$d = \operatorname{SEigen}(D^2 u)$$

Sochen, Kimmel,  
Malladi, 1998

$$\frac{\partial u}{\partial t} = \operatorname{div} \left( \frac{Du}{\sqrt{|Du|^2 + 1}} \right)$$

Note: Original Perona-Malik diffusion process is NOT anisotropic, although they erroneously said it was.




- A bilateral filter is an edge-preserving and noise reducing smoothing filter.
- Traditional filtering is **domain filtering**, and enforces closeness by weighing pixel values with weights that fall off with distance.
- Similarly, we define **range filtering**, which averages image values with weights that decay with dissimilarity.
- **Range filters**
  - are nonlinear because their weights depend on image intensity or color
  - preserve edges
  - by themselves, may distort an image's color map  
(see C. Tomasi et al., Bilateral Filtering for Gray and Color, ICCV 1998)
- **Bilateral filtering** combines range and domain filtering.
- The intensity value at each pixel in an image is replaced by a **weighted average of intensity values from nearby pixels**.
- This weight is based on a Gaussian distribution.  
The **weights** depend not only on Euclidean distance but also on the radiometric differences (e.g. color intensity).
- It preserves sharp edges by systematically looping through each pixel and attributing weights to the adjacent pixels accordingly.
- *OpenCV: Smooth()*

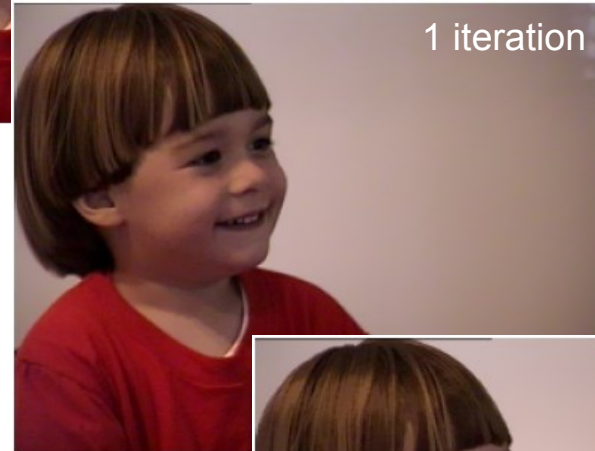
original



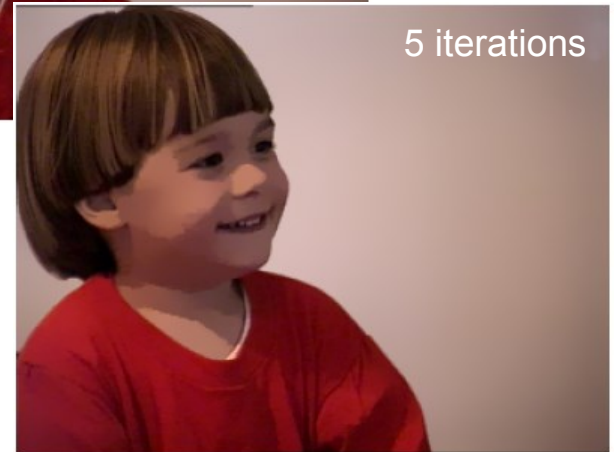
original



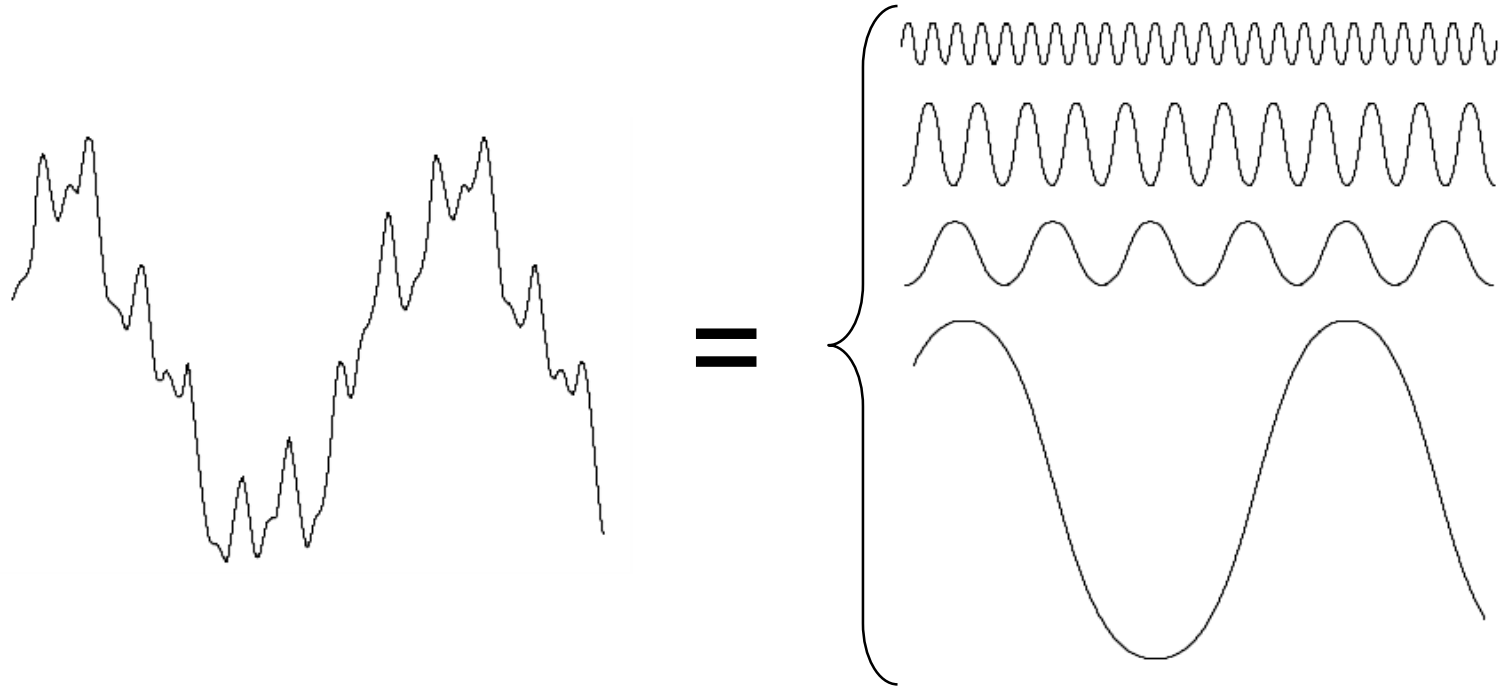
1 iteration



5 iterations



- An image can be represented in the frequency domain, using the Fourier Transform (FT).
- The FT encodes the amplitude and phase of each frequency component.
- The values near the origin of the transformed space are called low-frequency components of the FT, and those distant from the origin are the high-frequency components.
- Convolution in the image domain corresponds to multiplication in the spatial frequency domain.
- Therefore, convolution with large filters, which would normally be expensive in the image domain, can be implemented efficiently using the Fast Fourier Transform (FFT).
- This is an important technique in many image processing applications.



- Any function that periodically repeats itself can be expressed as a sum of sines and cosines of different frequencies each multiplied by a different coefficient  
– a *Fourier series*

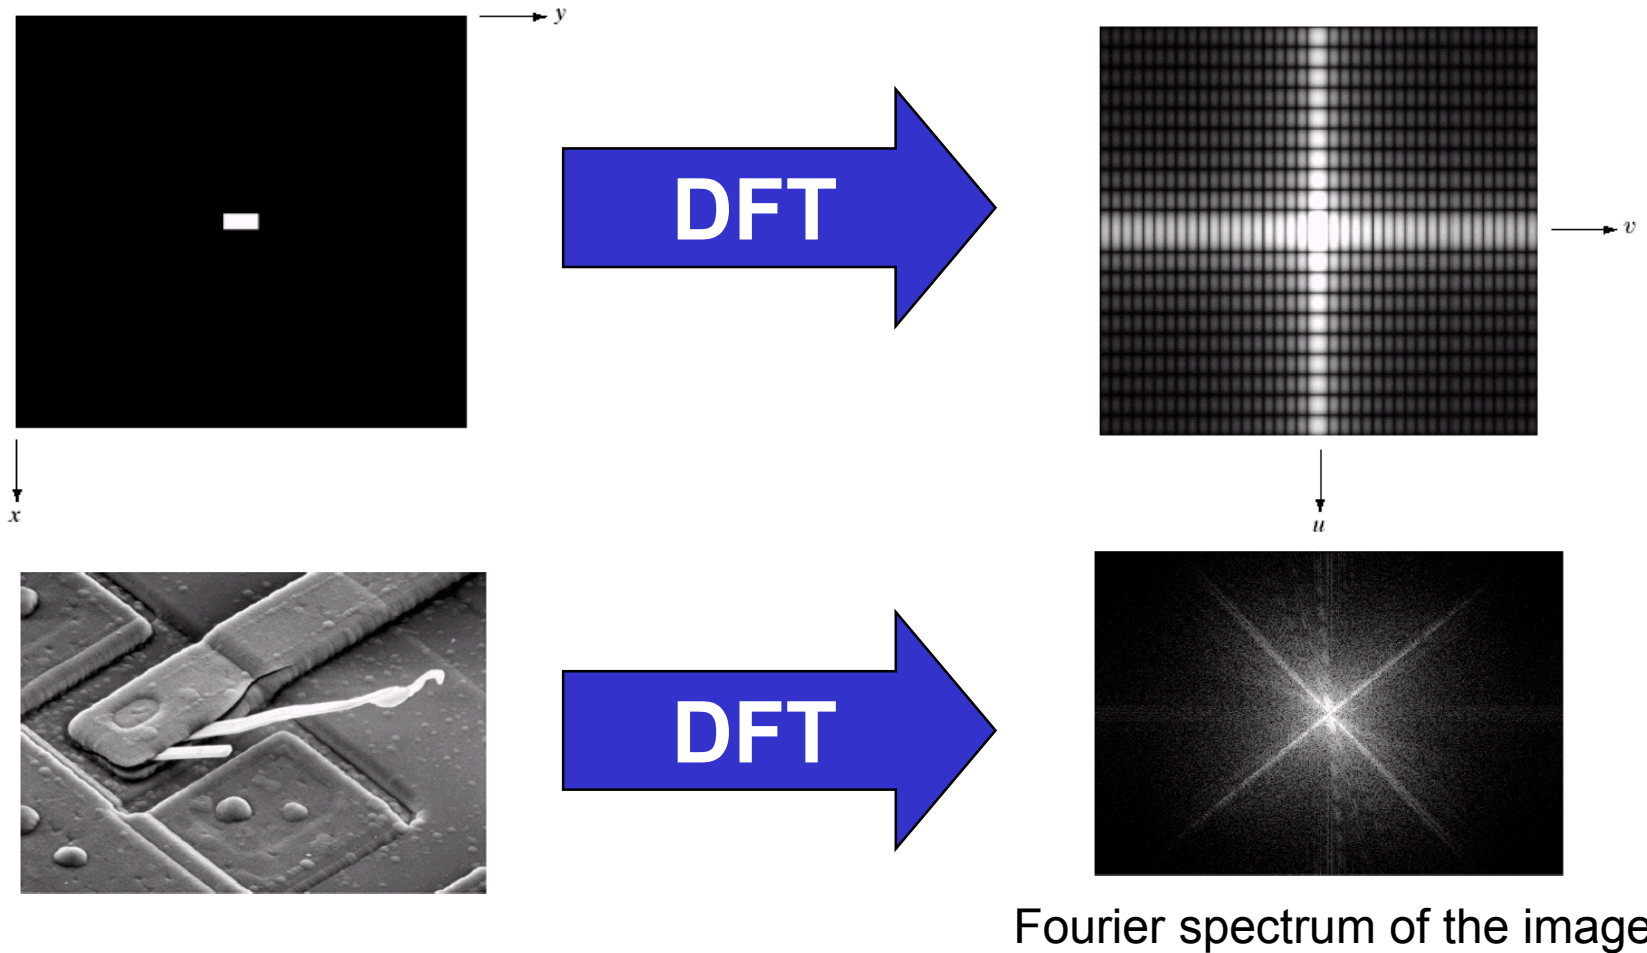
The *Discrete Fourier Transform* of  $f(x, y)$ ,  
for  $x = 0, 1, 2 \dots M-1$  and  $y = 0, 1, 2 \dots N-1$ ,  
denoted by  $F(u, v)$ , is given by the equation:

$$F(u, v) = \sum_{x=0}^{M-1} \sum_{y=0}^{N-1} f(x, y) e^{-j2\pi(ux/M + vy/N)}$$

for  $u = 0, 1, 2 \dots M-1$  and  $v = 0, 1, 2 \dots N-1$ .

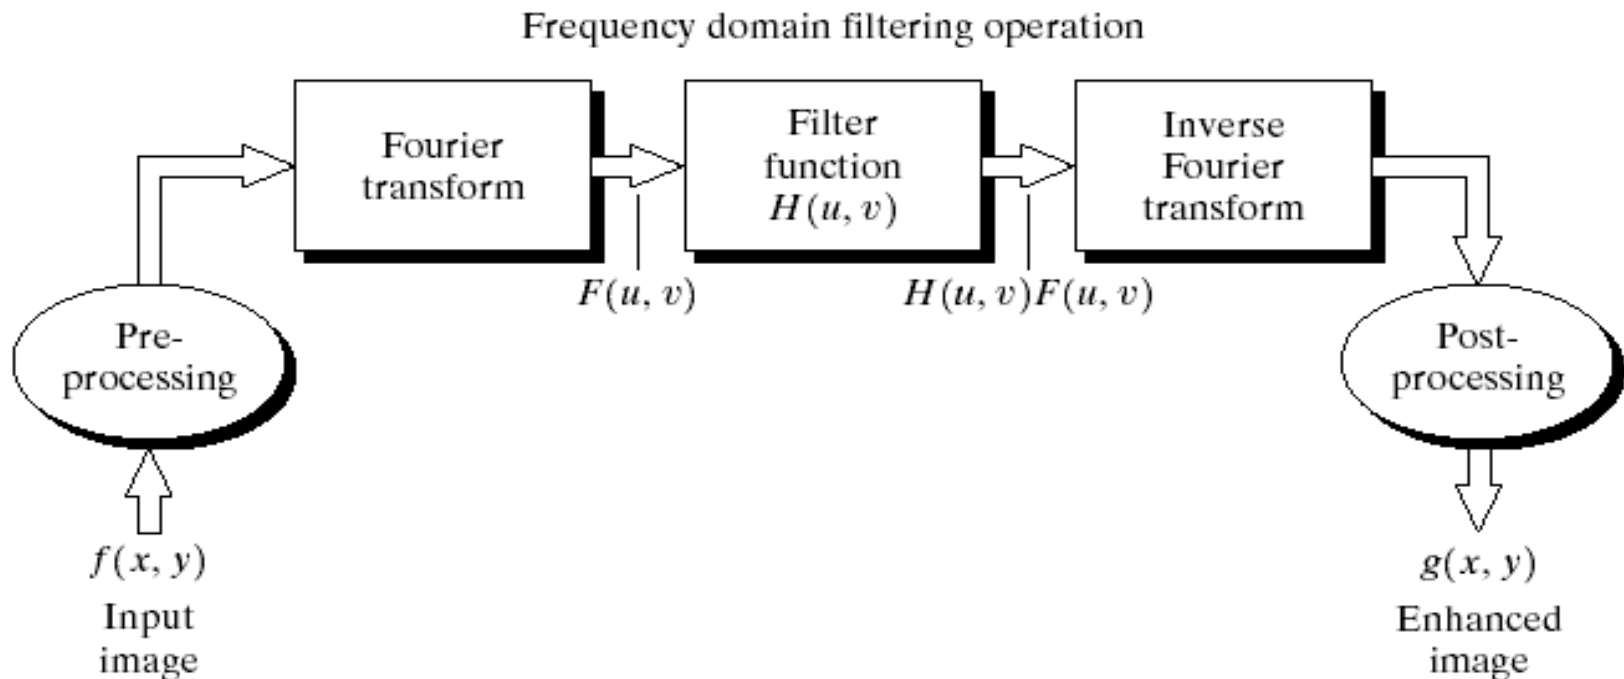


- The DFT of a 2D image can be visualised by showing the spectrum of the images component frequencies








To filter an image in the frequency domain:

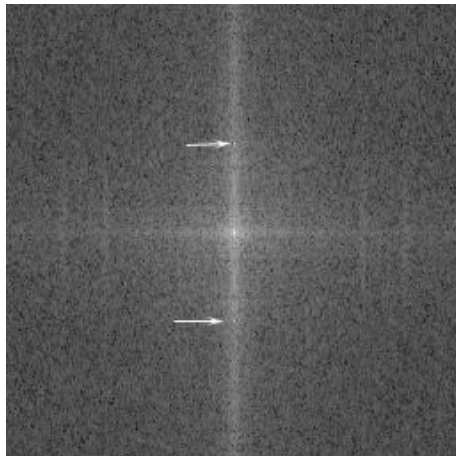
1. Compute  $F(u,v)$  the DFT of the image
2. Multiply  $F(u,v)$  by a filter function  $H(u,v)$
3. Compute the inverse DFT of the result



## Low Pass Filter



## High Pass Filter



After filtering

2D Fourier transform